

Unsupervised Learning


- Question: Which of the follow are unsupervised learning methods? (choose all that apply)
 - A. Principal components analysis
 - B. Causal Discovery
 - C. T-tests
 - D. GLMs
 - E. K-means
- All of the previous lectures have been on *supervised learning*
 - Supervised learning is the class of data method for mapping inputs to outputs, like prediction or inference
 - Supervised learning data: $(X_1, Y_1), \dots, (X_n, Y_n)$
 - Goal: Minimize some loss function: $L(f(X), Y)$
 - Typically, we think of X and Y as a joint distribution, $P(X, Y) = P(Y|X)P(X)$
 - In supervised learning, we are more interested $P(Y|X)$ (or $E(Y|X)$) than $P(X)$
- *Unsupervised learning* is a class of data methods for unlabeled data
 - Here, we are interested in learning properties of $P(X)$ without using labels, Y
 - examples: clustering, anomaly detection, latent variable analysis

Dimensionality Reduction


- Methods for reducing the number of variables in a dataset

Principle Components Analysis (PCA)

- Question: Have you ever used PCA?
 - A. Yes
 - B. No
- PCA tries to explain variance in the data using linear combinations of the variables



- [Developed in 1901 by Pearson \(https://en.wikipedia.org/wiki/Principal_component_analysis\)](https://en.wikipedia.org/wiki/Principal_component_analysis)



- PCA can be thought of as fitting a p -dimensional ellipsoid to the data, where each axis of the ellipsoid represents a principal component
- If some axis of the ellipsoid is small, then the variance along that axis is also small
- Assume X is an $n \times p$ *standardized* data matrix
 - standardized data: for each column, X_j of X , replace with $\frac{X_j - \bar{X}_j}{\sqrt{\text{var}(X_j)}}$
- First principal component: find unit-vector (direction), w , that best fits the data
 - find the directions within the data with the greatest variance
- Optimization:

$$w^* = \arg \max_{\|w\|=1} \widehat{\text{var}}(Xw) = \arg \max_{\|w\|=1} \|Xw\| = \arg \max_w \frac{w^T X^T X w}{w^T w}$$

where $\|w\| = \sqrt{\sum_i w_i^2}$ is the Euclidean norm

- Recall that $X^T X$ is the empirical covariance matrix
- Eigenvectors of $X^T X$ solve the problem
 - Lagrangian:

$$\mathcal{L}(w, \lambda) = w^T X^T X w - \lambda w^T w$$

- Gradient with respect to w :

$$\nabla_w \mathcal{L} = 2X^T X w - 2\lambda w = 0$$

- Then $X^T X w = \lambda w$, so principal components are eigenvectors of $X^T X$
- The first principal component has largest corresponding eigenvalue, λ
- The second principal component has second largest, etc

- Notice that there will be p total principal components
 - it is common to use the first k principal components with the that explain the majority of the variance, above some pre-set cutoff
- If w_1, w_2, \dots, w_k are the first $k < p$ components, we can regress an outcome on Xw_1, \dots, Xw_k in place of X
- All principal component vectors, w_1, \dots, w_p are mutually orthogonal (zero dot product or forming right angles)
- PCA is an old method that is still sometimes used but it has inspired many other methods

Iris Data

- 150 rows, 4 features
- Label (target): 0,1,2 (flower type)
- We only label for color coding but won't use in analysis
- Below is are scatter and density plots


```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
from sklearn.decomposition import PCA

from sklearn.datasets import load_iris
iris = load_iris()
label_data = pd.DataFrame(
    data= np.c_[iris['data'], iris['target']],
    columns= iris['feature_names'] + ['target'])
data = iris['data']

print(f'number of rows, number of columns: {data.shape}')

sns.pairplot(label_data, hue="target");
```

number of rows, number of columns: (150, 4)



Below: Using PCA to explain variance in data with `sklearn`

```
In [2]: pca = PCA()
pca.fit(data)
print(f'Variance explained in each component: \
{pca.explained_variance_ratio_}')
print(f'lambda values: {pca.singular_values_}')
print(f'All components: {pca.components_}')

Variance explained in each component: [0.92461872 0.05306648 0.01710
261 0.00521218]
lambda values: [25.09996044  6.01314738  3.41368064  1.88452351]
All components: [[ 0.36138659 -0.08452251  0.85667061  0.3582892 ]
 [ 0.65658877  0.73016143 -0.17337266 -0.07548102]
 [-0.58202985  0.59791083  0.07623608  0.54583143]
 [-0.31548719  0.3197231   0.47983899 -0.75365743]]
```


- Choosing the top two components
- Jointly they explain ~98% of the variance in features

```
In [3]: pca = PCA(n_components=2)
pca.fit(data)
pca_iris = pca.fit_transform(data)
pca_iris[0:5,]


Out[3]: array([[ -2.68412563,  0.31939725],
 [ -2.71414169, -0.17700123],
 [ -2.88899057, -0.14494943],
 [ -2.74534286, -0.31829898],
 [ -2.72871654,  0.32675451]])
```

Below: Labeled scatter plot of first two principal components

```
In [4]: pca_label = pd.DataFrame(  
    data= np.c_[pca_iris, iris['target']],  
    columns= ['PCA1', 'PCA2'] + ['target'])  
sns.scatterplot(data=pca_label, x='PCA1', y='PCA2', hue='target');
```



[Handwritten Digit Data \(https://archive.ics.uci.edu/ml/datasets/Pen-Based+Recognition+of+Handwritten+Digits\)](https://archive.ics.uci.edu/ml/datasets/Pen-Based+Recognition+of+Handwritten+Digits)



- 250 samples from 44 writers (according to UCI data repo)
- 1797 8x8 images transformed into feature vectors of length 64
- There is variation in writing style, thickness, and orientation


```
In [5]: from sklearn.datasets import load_digits
X, y = load_digits(return_X_y=True)
print(f'Dimensions of X: {X.shape}')
print(f'Dimensions of y: {y.shape}')
print('Looking at top left corner of data:')
X[0:5,0:10]
```

```
Dimensions of X: (1797, 64)
Dimensions of y: (1797,)
Looking at top left corner of data:
```

```
Out[5]: array([[ 0.,  0.,  5., 13.,  9.,  1.,  0.,  0.,  0.,  0.],
 [ 0.,  0.,  0., 12., 13.,  5.,  0.,  0.,  0.,  0.],
 [ 0.,  0.,  0.,  4., 15., 12.,  0.,  0.,  0.,  0.],
 [ 0.,  0.,  7., 15., 13.,  1.,  0.,  0.,  0.,  8.],
 [ 0.,  0.,  0.,  1., 11.,  0.,  0.,  0.,  0.,  0.]])
```

```
In [6]: pca = PCA()
pca.fit(X)
print(f'Variance explained in first 10 components:\n\
{np.round(pca.explained_variance_ratio_[0:10], 4)}')
print(f'Variances of first 10: \n\
{np.round(pca.singular_values_[0:10], 2)}')
```

```
Variance explained in first 10 components:
[0.1489 0.1362 0.1179 0.0841 0.0578 0.0492 0.0432 0.0366 0.0335 0.03
08]
Variances of first 10:
[567.01 542.25 504.63 426.12 353.34 325.82 305.26 281.16 269.07 257.
82]
```

- How well do the first 10 principal components predict the digit compared to the entire data?
- Below we'll use random forest on the full data and on the PCA transformed data
- Rather than using cross (or held out) validation, we will use out of bag sample accuracy

```
In [7]: pca10 = PCA(n_components=10)
pca10.fit(X)
X_pca10 = pca.fit_transform(X)

from sklearn.ensemble import RandomForestClassifier
rf_full = RandomForestClassifier(max_depth=2, random_state=0,
                                oob_score=True)

rf_full.fit(X, y)
print(f'Out of bag accuracy for full data: {rf_full.oob_score_}')

rf_pca = RandomForestClassifier(max_depth=2, random_state=0,
                                oob_score=True)


rf_pca.fit(X_pca10, y)
print(f'Out of bag accuracy PCA transformed data: \
{rf_pca.oob_score_}')
```

Out of bag accuracy for full data: 0.7835281023928771


Out of bag accuracy PCA transformed data: 0.7857540345019477

Independent Component Analysis (ICA)

- Question: Have you heard of ICA before?
 - A. Yes
 - B. No
- What if the sources of noise in the data are not orthogonal?
- Image below: left panel is PCA; right panel is ICA



- ICA attempts to factor a random matrix $X = AS$ where X is the $n \times p$ data matrix
- S is an $n \times p$ matrix called the latent source matrix
- A is a $p \times p$ matrix called the mixing matrix
- Goal: Estimate A so that the columns of S are *mutually independent*
- Interestingly, this is only possible for non-Gaussian distributions
- Cocktail Party Example: Assume that there are p microphones placed around a room with p people speaking at the same time
 - Goal: Use signal from each microphone to distinguish each speaker



- Without loss of generality, we can assume that $E(S) = 0$ and $\text{Cov}(S) = I_p$ (this covariance assume uncorrelated, not independent!)

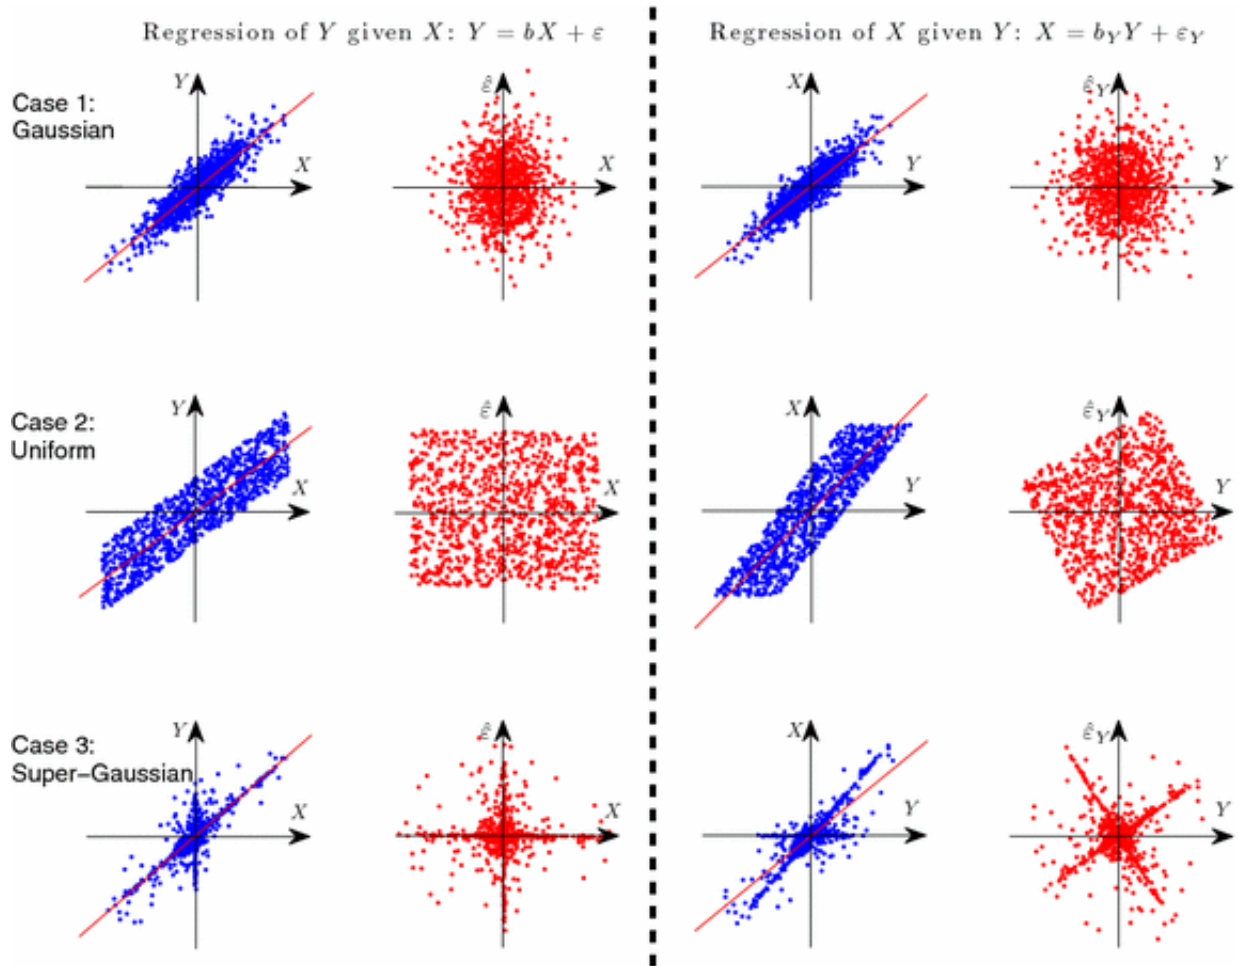
- Then

$$\text{Cov}(X) = \text{Cov}(AS) = AA^T$$

- If R is any orthogonal $p \times p$ matrix ($R^{-1} = R^T$), then

$$X = AS = AR^T RS = (AR^T)(RS)$$

- And, $\text{Cov}(RS) = RI_p R^T = I_p$, because correlation is not unique under rotations
- For Gaussian distribution, it is perfectly characterized by its first and second moments
 - For Gaussian distributions uncorrelated implied independent, which is why we assume non-Gaussian
- ICA uses third and higher moments to estimate A
 - In fact, it attempts to maximize divergence from Gaussianity
- ICA can be used for *causal discovery*
 - Causal discovery methods estimate a causal Bayesian network from data
 - Assume $X \rightarrow Y$ as $Y = bX + \epsilon$ where X and ϵ are independent random variables, and $E[\epsilon] = 0$ (zero-mean noise)
 - We only observe X and Y in the data but X and ϵ are the independent components
 - If X and ϵ are non-Gaussian, we can use ICA to determine causal ordering



Optimization

- There is a [lot of literature \(https://www.wiley.com/en-us/Independent+Component+Analysis-p-9780471405405\)](https://www.wiley.com/en-us/Independent+Component+Analysis-p-9780471405405) on ways to optimize ICA
- Using entropy and mutual information:
 - Recall entropy $H(Y) = - \int p(y) \log p(y) dy$ where $Y \sim p$
 - If $Y = (Y_1, \dots, Y_p)$, then the multivariate mutual information between all variables is

$$I(Y) = \sum_{j=1}^p H(Y_j) - H(Y)$$

- Recall that if $I(Y) = 0$ then all Y_j are mutually independent
- If $\text{Cov}(X) = I_p$ and W is a $p \times p$ orthogonal matrix, we can minimize

$$W^* = \arg \min_W I(WX) = \arg \min_W \sum_{j=1}^p H(w_j^T X) - H(X)$$

where w_j is the j th row of W

- Then use $A^* = (W^*)^{-1}$
- Unfortunately, estimating entropy can be challenging and sometimes computationally expensive
- FastICA:

- Interestingly, minimizing $I(WX)$ is equivalent to maximizing the departure of each $w_j^T X$ from Gaussianity
- Negentropy $J(Y_j) = H(Y_j) - H(Z_j)$ where Z_j is a Gaussian random variable with the same variance as Y_j (measure departure from Gaussianity)
- FastICA approximates negentropy with

$$J(w_j^T) \approx (E[G(w_j^T X)] - E[G(w_j^T Z)])$$

where $G(y) = \frac{1}{a} \log \cosh y$ for $a \in (1, 2)$

ICA notes based on these [slides \(http://statweb.stanford.edu/~tibs/sta306bfiles/ica.pdf\)](http://statweb.stanford.edu/~tibs/sta306bfiles/ica.pdf)

- The code below uses FastICA to transform the the data
- Again, we use random forest to see how well the transformed data predicts the label
- Note: it can be difficult to optimize ICA

```
In [8]: from sklearn.decomposition import FastICA
transformer10 = FastICA(n_components=10, random_state=0)
X_ica10 = transformer10.fit_transform(X)


rf_ica = RandomForestClassifier(max_depth=2, random_state=0,
                               oob_score=True)

rf_ica.fit(X_ica10, y)
print(f'Out of bag accuracy for ICA transformed data: \
{rf_full.oob_score_}')
```

Out of bag accuracy for ICA transformed data: 0.7835281023928771

Cluster Analysis (clustering)

- Clustering attempts to group together observations that are close to each other
- This can take different forms
- It might be that data points are grouped around central points



- In other cases, groups are formed by proximity to other points in the group, like a chain



K-means

- Goal: group points using proximity to K unknown cluster centers
- This approach would work better for data in the top image, but not the bottom image
- Assume X_1, \dots, X_n is a sample where $X_i \in \mathbb{R}^d$ for $i = 1, \dots, n$
- Let $C = \{c_1, \dots, c_K\}$, $c_k \in \mathbb{R}^d$, be variable point that we will use to find cluster centers

- [Lloyd's algorithm](https://en.wikipedia.org/wiki/Lloyd%27s_algorithm) (https://en.wikipedia.org/wiki/Lloyd%27s_algorithm) has two steps:

1. Assign a cluster: for each point, $i = 1, \dots, n$, find the closest $c_k \in C$,

$$z_i = \arg \min_{1 \leq k \leq K} \|X_i - c_k\|$$

2. Recenter cluster: update each c_k with the new group,


$$c_k \leftarrow \frac{1}{n_k} \sum_{i: z_i = k} X_i$$

$$n_k = \#\{i : z_i = k\}$$

- We can initialize this algorithm with K random points in \mathbb{R}^d
 - These two steps are repeated until C , the cluster center points, converge (until they stop moving between steps)
- Optimization:

$$\frac{1}{n} \sum_{i=1}^n \min_{1 \leq k \leq K} \|X_i - c_{z_i}\|^2$$

- Once the cluster centers are set, they create a *Voronoi Tessellation*:



- Below: we run this algorithm on Iris PCA-transformed data

```
In [9]: from sklearn.cluster import KMeans
kmeans = KMeans(n_clusters=3, random_state=0).fit(pca_iris)


# switching labels to match
def switch_1_2(x):
    if x==0.0:
        return 0.0
    if x==1.0:
        return 2.0
    if x==2.0:
        return 1.0
kmeans_labs = np.vectorize(switch_1_2)(kmeans.labels_)

pca_label = pd.DataFrame(
    data= np.c_[pca_iris, iris['target'], kmeans_labs],
    columns= ['PCA1', 'PCA2', 'True Label', 'K means'])
def class_error(row):
    return row['True Label'] == row['K means']
pca_label['Cluster'] = pca_label.apply(class_error, axis=1)


sns.scatterplot(data=pca_label, x='PCA1', y='PCA2', hue='K means',
                style='Cluster', style_order=[True,False]);
centers = pd.DataFrame(
    kmeans.cluster_centers_, columns=['PCA1', 'PCA2'])
sns.scatterplot(data=centers, x='PCA1', y='PCA2', color='red');
print(f'Number of misclassifications:')
pca_label.Cluster.value_counts()
```

Number of misclassifications:

```
Out[9]: True      133
        False     17
        Name: Cluster, dtype: int64
```




- Above: Showing K means groups with error as 'x' point markers
- The red points show the cluster centers
- This model does not account for different sized cluster



Gaussian Mixture Model and EM Clustering

- Another way to think about clustering is with a *Gaussian mixture model* (https://en.wikipedia.org/wiki/Mixture_model) (GMM)



- Assume that our data points are each generated from different Gaussian distributions
- All we see is the location, not the *latent* variable which determines which distribution each sample was generated from
- x_1, \dots, x_n is a random sample
- Each x_i is associated with a latent random variable, z_i , that indicates the cluster/distribution of x_i
- Density of z is multinomial with probabilities, π_1, \dots, π_K
- Let (μ_k, Σ_k) for $j = 1, \dots, K$ be the mean vector and covariance matrix for K multivariate Gaussian distributions
- In a Gaussian mixture model, each data point is generated from one of the K multivariate Gaussian distribution,

$$x_i \sim \text{Gaussian}(\mu_k, \Sigma_k)$$

with probability π_k

- For $i = 1, \dots, n$, if $z_i = k$ then

$$x_i \sim \text{Gaussian}(\mu_k, \Sigma_k) := p(x_i | z_i = k)$$

- Joint distribution of x and z

$$p(x, z) = \prod_{k=1}^K \pi_k^{I(z=k)} \cdot [\text{Gaussian}(x; \mu_k, \Sigma_k)]^{I(z=k)}$$


- Marginal distribution x (summing over the values of z)

$$p(x) = \sum_{k=1}^K \pi_k \cdot \text{Gaussian}(x; \mu_k, \Sigma_k)$$

- Each point is generated from one of these Gaussian distributions but we won't know which


```
In [10]: dist_probs = [0.2, 0.3, 0.5]
data = []
samp_size = 200

for _ in range(samp_size):
    k = np.random.choice(np.arange(0, 3), p=dist_probs)
    if k==0:
        obs = np.random.multivariate_normal(
            [-2,1], [[1,0.2], [0.2,1]]).tolist()
    elif k==1:
        obs = np.random.multivariate_normal(
            [0,-2], [[1,0], [0,1]]).tolist()
    else:
        obs = np.random.multivariate_normal(
            [1,1], [[1,0.3], [0.3,1]]).tolist()
    obs.append(k)
    data.append(tuple(obs))
df = pd.DataFrame(data, columns=['x1', 'x2', 'label'])
sns.scatterplot(x='x1', y='x2', data=df);
```



Below: color by label

```
In [11]: sns.scatterplot(x='X1', y='X2', data=df, hue='label');
```



EM Algorithm

- EM stands for Expectation Maximization
- We assume that the sample is generated from a Gaussian mixture model
- Using the Gaussian mixture likelihood, cycle between
 - Imputing the expected cluster (probability) of each point, $i = 1, \dots, n$
 - Recalculate GMM parameters, π_k, μ_k, Σ_k for $k = 1, \dots, K$, to maximize the likelihood
- Cosma Shalizi: "one man's vicious circle is another man's successive approximation procedure"
- Sample marginal likelihood of x :

$$L(\pi, \mu, \Sigma) = \prod_{i=1}^n \sum_{k=1}^K \pi_k \cdot \text{Gaussian}(x_i; \mu_k, \Sigma_k)$$

where $\pi = (\pi_1, \dots, \pi_K)$, $\mu = (\mu_1, \dots, \mu_K)$, $\Sigma = (\Sigma_1, \dots, \Sigma_K)$

- **E step:** Calculate expected value of cluster assignments for each point, x_i , using the values of π, μ, Σ from last step,

$$p^{\text{new}}(z_i = k) \leftarrow E[z_i = k | \pi, \mu, \Sigma]$$

- This expected value is really $P(z_i = k)$ (expected value of an indicator function is a probability)
- Calculate these probabilities as

$$p^{\text{new}}(z_i = k) \leftarrow \frac{\pi_k^{\text{old}} \cdot \text{Gaussian}(x_i; \mu_k^{\text{old}}, \Sigma_k^{\text{old}})}{\sum_{j=1}^K \pi_j^{\text{old}} \cdot \text{Gaussian}(x_i; \mu_j^{\text{old}}, \Sigma_j^{\text{old}})}$$

- To make the notation in the next step clearer, let

$$r_i(k) \leftarrow p^{\text{new}}(z_i = k)$$

- For each point, we have a vector of assignment probabilities

$$[r_i(1), \dots, r_i(K)]$$

- We can think about this vector as the *responsibility* that the GMM takes for explaining observation i
- **M step:** Calculate maximum likelihood estimates for all parameters, π, μ, Σ using the expected cluster assignments, $[r_i(1), \dots, r_i(K)]$
 - Goal: optimize likelihood of joint distribution $p(x, z)$

$$\pi^{\text{new}}, \mu^{\text{new}}, \Sigma^{\text{new}} \leftarrow \arg \max_{\pi, \mu, \Sigma} E \left[\log \left(\prod_{i=1}^n p(x_i, z_i) \right) \right]$$

- This can be hard because L is not convex with respect to π, μ, Σ
- Conditional density of x_i given z_i :

$$p(x_i | z_i) = \prod_{k=1}^K [\text{Gaussian}(x_i; \mu_k, \Sigma_k)]^{I(z_i=k)}$$

- Density of z_i :

$$p(z_i) = \prod_{k=1}^K \pi_k^{I(z_i=k)}$$

- Likelihood:

$$\begin{aligned} & E \left[\log \left(\prod_{i=1}^n p(x_i, z_i) \right) \right] \\ &= E \left[\log \left(\prod_{i=1}^n p(z_i) p(x_i | z_i) \right) \right] \\ &= E \left[\log \left(\prod_{i=1}^n \prod_{k=1}^K \pi_k^{I(z_i=k)} \cdot [\text{Gaussian}(x_i; \mu_k, \Sigma_k)]^{I(z_i=k)} \right) \right] \\ &= E \left[\sum_{i=1}^n \sum_{k=1}^K I(z_i = k) \log \pi_k + I(z_i = k) \log [\text{Gaussian}(x_i; \mu_k, \Sigma_k)] \right] \\ &= \sum_{i=1}^n \sum_{k=1}^K r_i(k) \log \pi_k + r_i(k) \log [\text{Gaussian}(x_i; \mu_k, \Sigma_k)] \end{aligned}$$

- This can be solved using Lagrange multiplier with the constraint that $\sum_{k=1}^K \pi_k = 1$

$$\mathcal{L} = \sum_{i=1}^n \sum_{k=1}^K r_i(k) \log \pi_k + r_i(k) \log [\text{Gaussian}(x_i; \mu_k, \Sigma_k)] - \lambda \left(\sum_{k=1}^K \pi_k - 1 \right)$$

- Differentiate with respect to parameters, π, μ, Σ and λ then set equal to zero
- Solutions:

$$\begin{aligned}
 n_k^{\text{new}} &\leftarrow \sum_{i=1}^n r_i(k) \\
 \pi_k^{\text{new}} &\leftarrow \frac{n_k^{\text{new}}}{n} \\
 \mu_k^{\text{new}} &\leftarrow \frac{1}{n_k^{\text{new}}} \sum_{i=1}^n r_i(k) x_i \\
 \Sigma_k^{\text{new}} &\leftarrow \frac{1}{n_k} \sum_{i=1}^n r_i(k) (x_i - \mu_k^{\text{new}})(x_i - \mu_k^{\text{new}})^T
 \end{aligned}$$

- The algorithm must begin initialized with parameters, π, μ, Σ
- The algorithm cycles through these steps until the parameters converge

```

In [12]: from sklearn.mixture import GaussianMixture

# n_components here is number of clusters
gm = GaussianMixture(n_components=3, random_state=0).fit(pca_iris)

# switching cluster labels
gm_labs = np.vectorize(switch_1_2)(gm.fit_predict(pca_iris))

pca_label_gm = pd.DataFrame(
    data= np.c_[pca_iris, iris['target'], gm_labs],
    columns= ['PCA1', 'PCA2', 'True Label', 'GM'])
def class_error(row):
    return row['True Label'] == row['GM']
pca_label_gm['Cluster'] = pca_label_gm.apply(class_error, axis=1)

sns.scatterplot(data=pca_label_gm, x='PCA1', y='PCA2', hue='GM',
                style='Cluster', style_order=[True, False]);
centers = pd.DataFrame(
    gm.means_, columns=['PCA1', 'PCA2'])
sns.scatterplot(data=centers, x='PCA1', y='PCA2', color='red');
print(f'Number of misclassifications:')
pca_label_gm.Cluster.value_counts()

```

Number of misclassifications:

```
Out[12]: True      147  
False       3  
Name: Cluster, dtype: int64
```

